

# Linux Load Averages: Solving the Mystery

## Load Average in LINUX

Autour: M.M.Mirzaei



# Table of contents

**Basics of Load Average in Linux**

**01**

**CPU Load Average Calculation**

**02**

**The Problem Appear**

**03**

**The origin of uninterruptible**

**04**

fi..	finish_task	finish_task
sc..	schedule	schedule
sc..	schedule_task	schedule_task
wa..	wait_for_completion	wait_for...
_r..	__wait_rcu	__wait_rcu
rc..	synchronization	synchroniz...
re..	memcg_dealloc	memcg_deal...
me..	memcg_offline	memcg_offli...
me..	mem_cgroup	mem_cgrou...
cs..	css_killed	css_kille...
s..	process_on_rq	process_o...
r_t..	worker_thread	worker_th...
d	kthread	kthread
m..	ret_from_f	ret_from...
er/	kworker/6:2	kworker/8:2

What is  
**LOAD  
AVERAGE**  
in Linux



## Introduction

---

Load Average in Linux is a metric that is used by Linux users to keep track of system resources. It also helps you monitor how the system resources are engaged.



0 1

# Basics of Load Average in Linux



# Load? Why Load Average?

## Theory

---

It refers to the number of processes which are either currently being executed by the CPU or are waiting for execution.

An idle system has a load of 0. With each process that is being executed or is on the waitlist, the load increases by one.

On its own, **the load doesn't give any useful information** to the user. **The load can change in split seconds.** This is because the number of processes using or waiting for the CPU time doesn't remain constant. This is why we use Load Average in Linux to monitor resource usage.

# How to Check the Load Average in Linux

sudo apt-get install glances

## Terminal Command

glances

top

uptime

```
ubuntu (Ubuntu 18.04 64bit / Linux 4.15.0-96-generic) - IP: 108  
12/24 Pub: 108 12 Uptime: 2 days, 16:30:56  
AWS EC2 instance 1340aa
```

```
CPU [ 2.9% ] CPU 2.9% nice: 0.0% ctx_sw: 246 MEM 59.8% active: 283M SWAP 0.0%  
MEM [ 59.8% ] user: 2.3% irq: 0.0% inter: 75 total: 481M inactive: 88.5M total: 0 1 min: 0.14  
SWAP [ 0.0% ] system: 0.7% iowait: 0.0% sw_int: 123 used: 288M buffers: 10.1M used: 0 5 min: 0.12  
idle: 97.1% steal: 0.0% cached: 195M free: 0 15 min: 0.05
```

```
LOAD 1-core
```

```
root@ubuntu:~# uptime
```

```
12:28:35 up 2 days, 16:26, 1 user, load average: 0.03, 0.03, 0.01  
root@ubuntu:~#
```

```
top - 12:31:47 up 2 days, 16:29, 1 user, load average: 0.34, 0.14, 0.05
```

```
Tasks: 112 total, 2 running, 72 sleeping, 0 stopped, 0 zombie
```

```
%Cpu(s): 1.0 us, 1.0 sy, 0.0 ni, 98.0 id, 0.0 wa, 0.0 hi, 0.0 si, 0.0 st
```

```
KiB Mem : 492632 total, 68960 free, 238684 used, 184988 buff/cache
```

```
KiB Swap: 0 total, 0 free, 0 used. 228020 avail Mem
```

PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+	COMMAND
-----	------	----	----	------	-----	-----	---	------	------	-------	---------

19305	root	20	0	44392	3884	3268	R	0.3	0.8	0:00.02	top
-------	------	----	---	-------	------	------	---	-----	-----	---------	-----

19306	root	20	0	114248	6952	5912	S	0.3	1.4	0:00.01	sshd
-------	------	----	---	--------	------	------	---	-----	-----	---------	------

19308	root	20	0	114248	7616	6024	S	0.3	1.4	0:00.01	sshd
-------	------	----	---	--------	------	------	---	-----	-----	---------	------

# What those three numbers really mean?

- The first value depicts the average load on the CPU for the last minute.
- The second gives us the average load for the last 5-minute interval
- The third value gives us the 15-minute average load

```
root@ubuntu:~# uptime  
12:28:35 up 2 days, 16:26, 1 user, load average: 0.03, 0.03, 0.01  
root@ubuntu:~# █
```

This helps a user get an idea of how the CPU is being utilized by the processes on a system over time.



# How to Check the Load Average in Linux



Linus Torvalds  
torvalds

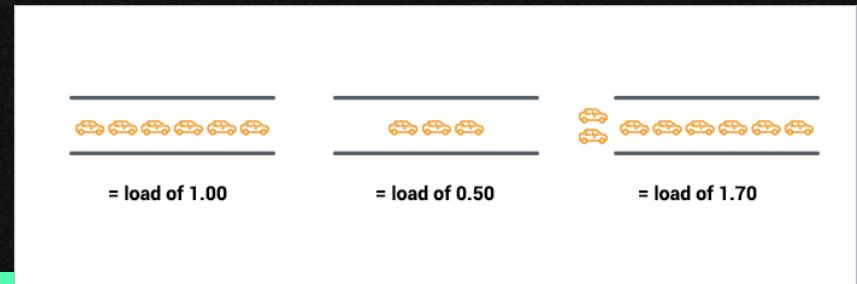


```
1 // SPDX-License-Identifier: GPL-2.0
2 /*
3  * kernel/sched/loadavg.c
4  *
5  * This file contains the magic bits required to compute the global loadavg
6  * figure. Its a silly number but people think its important. We go through
7  * great pains to make it work on big machines and tickless kernels.
8  */
9 #include <linux/sched/nohz.h>
10 #include "sched.h"
11
12 /*
13  * Global load-average calculations
14  *
15  * We take a distributed and async approach to calculating the global load-avg
16  * in order to minimize overhead.
17  *
18  * The global load average is an exponentially decaying average of nr_running +
19  * nr_uninterruptible.
20  *
21  * Once every LOAD_FREQ:
```

```
$ cat /proc/loadavg
25.72 23.19 23.35 42/3411 43603
```

# When should you be worried?

- 0.00 means there's no traffic on the bridge at all.
- 1.00 means the bridge is *exactly* at capacity.
- Over 1.00 means there's backup.



02



# CPU Load Average Calculation

---



F[x]

$$\text{new\_load} = \text{old\_load} \times e^{-\frac{\Delta t}{\tau}} + \text{current\_load} \times \left(1 - e^{-\frac{\Delta t}{\tau}}\right)$$

—

*Exponential Moving Average (EMA)*

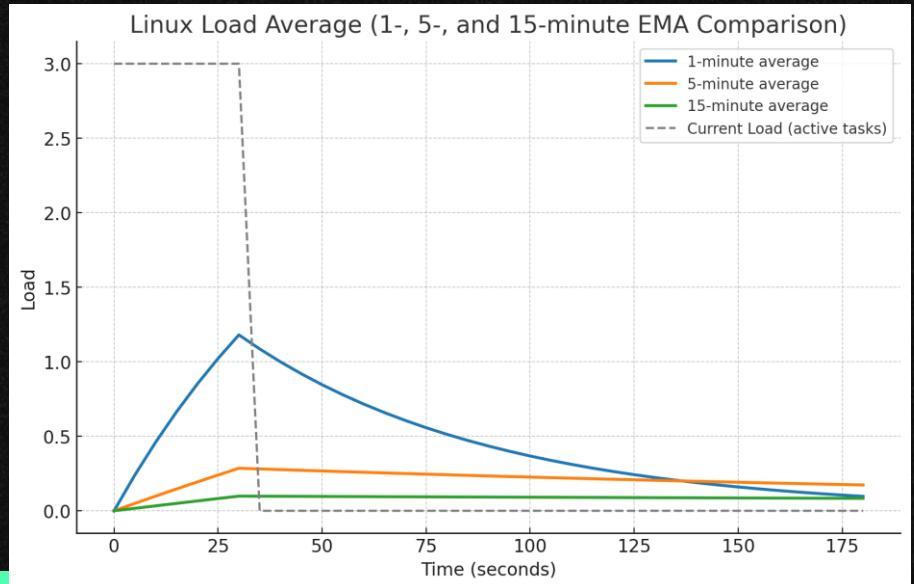
Time Constant

Decay Factor

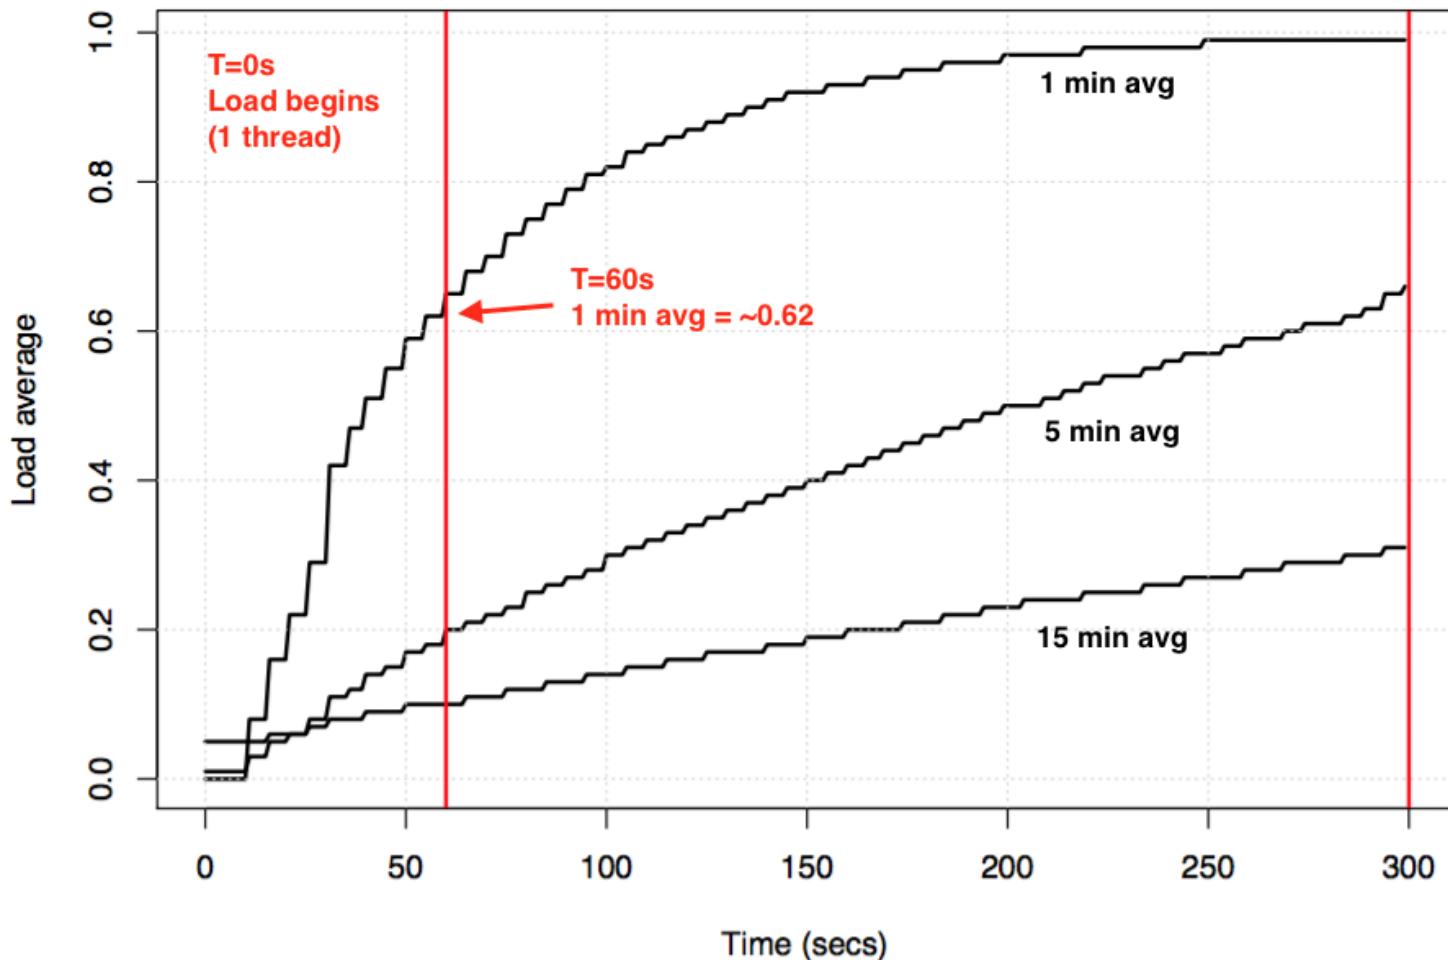


# Why Uses Exponential Moving Average (EMA)?

- Prevents rapid fluctuations in displayed load values.
- Provides a better “**sense of system pressure**” instead of raw, noisy data.
- Helps visualize **long-term patterns** of CPU and I/O demand.



## Load averages: 1, 5, 15 min



# Where it comes from?

Inside the Linux kernel source  
(kernel/sched/loadavg.c)

```
#define LOAD_FREQ (5*HZ) /* 5 sec intervals */  
#define EXP_1 1884 /* 1/exp(5sec/1min) */  
#define EXP_5 2014 /* 1/exp(5sec/5min) */  
#define EXP_15 2037 /* 1/exp(5sec/15min) */
```

Average	$\tau$ (s)	$\exp(-5/\tau)$	In decimal	In kernel scaled form
1 min	60	$e^{-5/60}$	0.9200	1884
5 min	300	$e^{-5/300}$	0.9835	2014
15 min	900	$e^{-5/900}$	0.9945	2037

03

# The Problem Appear



# The problem appear



## What's the difference between load average and CPU load?

Asked 14 years, 11 months ago Modified 14 years, 11 months ago Viewed 104k times

Here's the output of top:

41 top - 23:30:49 up 2:18, 1 user, load average: 4.36, 4.36, 4.39  
Tasks: 105 total, 2 running, 103 sleeping, 0 stopped, 0 zombie  
Cpu(s): 3.6%us, 8.0%sy, 0.0%ni, 73.4%id, 9.0%hi, 4.9%si, 0.0%st  
Mem: 2029820k total, 1979312k used, 50508k free, 6828k buffers  
Swap: 5947384k total, 0k used, 5947384k free, 1855304k cached

How can the load average be high while the CPU load is low.

How is load average calculated?

cpu-load

Share Improve this question Follow

asked Oct 25, 2010 at 21:33



Georg Scholly  
3,751 ● 5 ● 27 ● 34

What are you running, it seems your system is working on something. I get loads like that when I compress a bunch of stuff while compiling. – [NightwishFan](#) Oct 25, 2010 at 22:15

I always thought of the load as the "number of processes waiting to be executed", but I'm not really sure if that is correct. – [pableu](#) Oct 26, 2010 at 7:21

I have 6 rackspace servers. It is currently telling me that I had an extended period (an hour) with a load of over 20 on a 2 processor machine. I am scratching my head about CPU load too. – [MagicLAMP](#) Jan 25, 2018 at 23:48

Add a comment

2 Answers

Sorted by: Highest score (default)

[This site](#) does a good job of explaining it. Basically, load average is the amount of traffic to your CPU(s) over the past 1, 5, and 15 minutes. Generally you want this number to be below the number of CPU(s)/cores you have. 1.0 on a single core machine means it's using the CPU to its maximum, and anything above that means things are getting queued.

40

The CPU line in your top output is the current usage broken down by process types.

Share Improve this answer Follow

answered Oct 25, 2010 at 22:05



Daniel R  
517 ● 3 ● 5

10 As a note, load average also includes processes waiting on I/O: [serverfault.com/a/524818/27813](#)  
– [rogerdpark](#) Jul 19, 2013 at 17:49

1 The linked site has a really good explanation. – [asheeshr](#) Sep 22, 2013 at 10:53

Add a comment

What Inigoesdr and the site he/she points to write is more or less correct, but remember that the "load average" isn't really a "regular" mathematical average, it's a *exponentially damped/weighted moving average*.

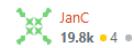
16

[This](#) is a very good and in-depth article on the topic of CPU percentage and load average, and how they are calculated in linux. Wikipedia also has a good [article](#) on it (explaining some differences between load average on linux vs. most UNIX systems for example).

Share Improve this answer Follow

edited Oct 26, 2010 at 0:01

answered Oct 25, 2010 at 23:55



JanC  
19.8k ● 4 ● 46 ● 50

Add a comment

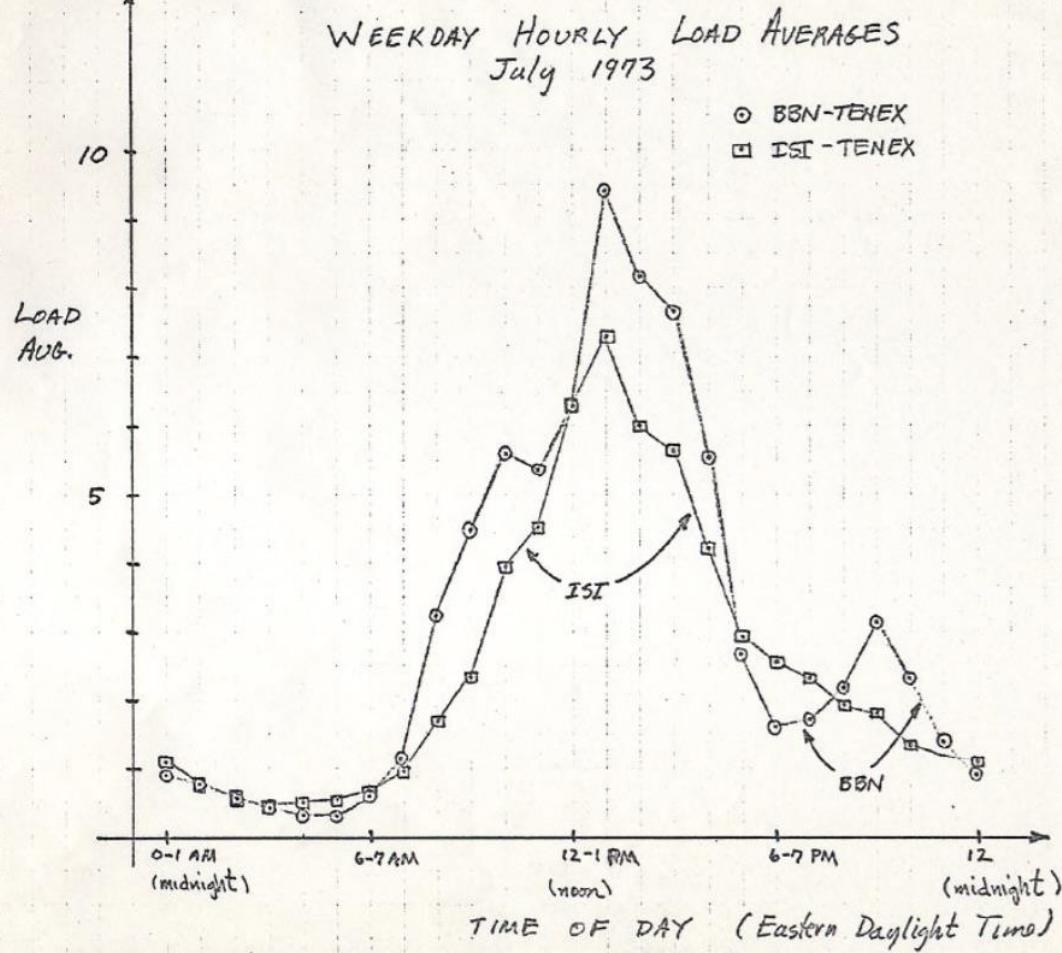
# History

There's a nice description of this in [RFC 546](#) titled "TENEX Load Averages", August 1973:

RFC 546      TENEX Load Averages for July 1973      10 August 1973

Endnotes

[1] The TENEX load average is a measure of CPU demand. The load average is an average of the number of runnable processes over a given time period. For example, an hourly load average of 10 would mean that (for a single CPU system) at any time during that hour one could expect to see 1 process running and 9 others ready to run (i.e., not blocked for I/O) waiting for the CPU.

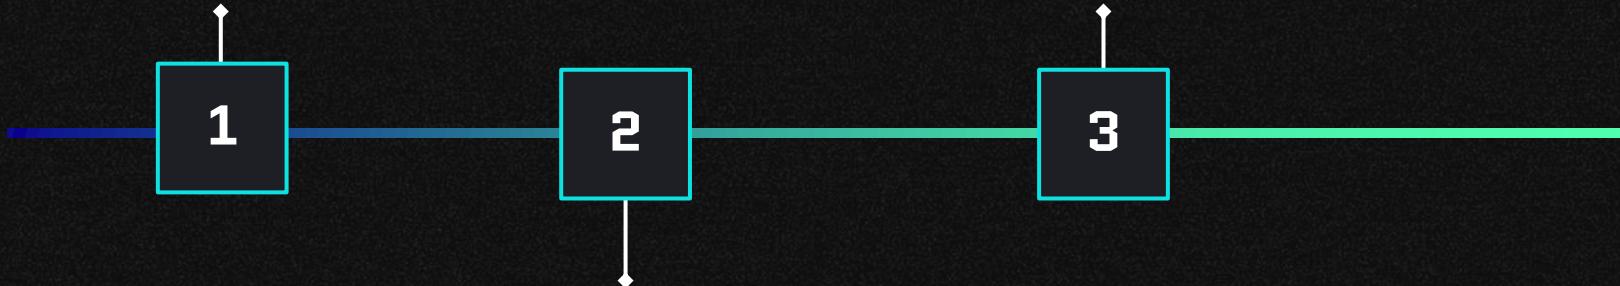




# TimeLine

## 1973 TENEX

Process that currently being executed by the CPU or are waiting for execution.



## 1993

The load averages were deliberately changed to reflect demand for other system resources, not just CPUs. Linux changed from "CPU load averages" to what one might call "system load averages".



04



# The origin of uninterruptible

# How and where they change it?

Compressed Mailbox File From 1993 On Oldlinux.Org

```
From: Matthias Urlichs <urlichhs@smurf.sub.org>
Subject: Load average broken ?
Date: Fri, 29 Oct 1993 11:37:23 +0200

The kernel only counts "runnable" processes when computing the load average.
I don't like that; the problem is that processes which are swapping or
waiting on "fast", i.e. noninterruptible, I/O, also consume resources.

It seems somewhat nonintuitive that the load average goes down when you
replace your fast swap disk with a slow swap disk...

Anyway, the following patch seems to make the load average much more
consistent wrt the subjective speed of the system. And, most important, the
load is still zero when nobody is doing anything. ;-)

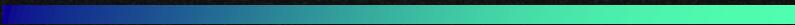
--- kernel/sched.c.orig Fri Oct 29 10:31:11 1993
+++ kernel/sched.c Fri Oct 29 10:32:51 1993
@@ -414,7 +414,9 @@
     unsigned long nr = 0;

     for(p = &LAST_TASK; p > &FIRST_TASK; --p)
-        if (*p && (*p)->state == TASK_RUNNING)
+        if (*p && ((*p)->state == TASK_RUNNING) ||
+            (*p)->state == TASK_UNINTERRUPTIBLE) ||
+            (*p)->state == TASK_SWAPPING))
         nr += FIXED_1;
     return nr;
 }
-- 
Matthias Urlichs      \ XLink-POP N|rnberg | Email: urlichhs@smurf.sub.org
Schleiermacherstra_e 12 \ Unix+Linux+Mac | Phone: ...please use email.
90491 N|rnberg (Germany) \ Consulting+Networking+Programming+etc'ing    42
```

# The Linux Process States

- Running or Runnable (R)
- Uninterruptible Sleep (D)
- Interruptable Sleep (S)
- Stopped (T)
- Zombie (Z)

**R → running ↔ S/D → T (optional) → Z → gone.**



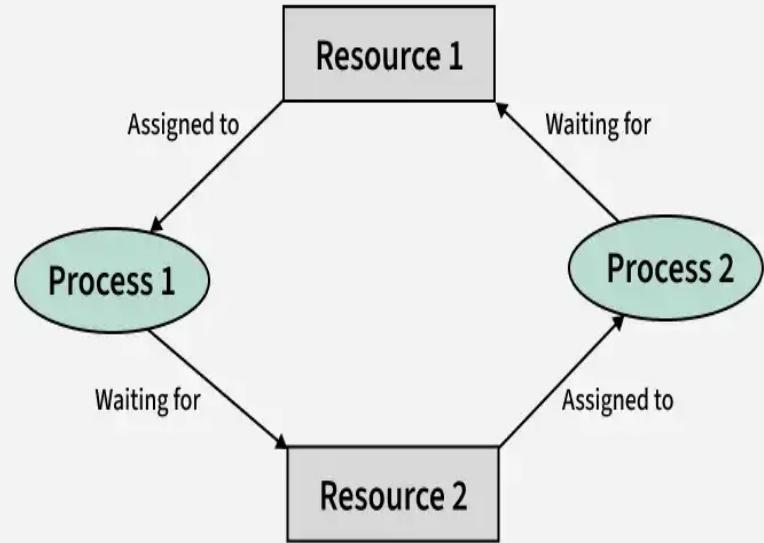
# Add Uninterruptible Tasks

When load averages first appeared in Linux, they reflected CPU demand, as with other operating systems. But later on Linux changed them to include not only runnable tasks, but also tasks in the uninterruptible state (TASK\_UNINTERRUPTIBLE or nr\_uninterruptible).

Adding the uninterruptible state means that Linux load averages can increase due to a disk (or NFS) I/O workload, not just CPU demand.

# Causes Of Uninterruptible Tasks [D State]

- Disk I/O (read/write to a block device)
- NFS or network storage operation
- A hardware controller response
- **Kernel lock or driver issue**





# Resources

## Context:

- ❑ <https://www.brendangregg.com/blog/2017-08-08/linux-load-averages.html>
- ❑ <https://www.digitalocean.com/community/tutorials/load-average-in-linux>
- ❑ <https://www.scoutapm.com/blog/understanding-load-averages>
- ❑ <https://superuser.com/questions/1402079/understanding-load-average-and-cpu-in-linux>
- ❑ <https://github.com/torvalds/linux/blob/master/kernel/sched/loadavg.c>
- ❑ <https://datatracker.ietf.org/doc/html/rfc546>

## Photos:

- <https://ahelpme.com/software/grafana/create-graph-for-linux-load-average-using-grafana-influxdb-and-collectd/>
- <https://www.brendangregg.com/blog/images/2017/rfc546.jpg>
- <https://www.geeksforgeeks.org/operating-systems/introduction-of-deadlock-in-operating-system/>
- <https://github.com/torvalds>

